

MLP_classifier_parameter

September 15, 2025

```
[17]: import ipynbname
import torch

if torch.cuda.is_available():
    device = torch.device("cuda")
elif torch.backends.mps.is_available():
    device = torch.device("mps")
else:
    device = torch.device("cpu")

print(f"Using device: {device}")
```

Using device: mps

```
[18]: %load_ext autoreload
%autoreload 2
from libraries.embeddings_utils import *
from libraries.classifier_utils import *
from libraries.file_manager_utils import *
from libraries.EmbeddingNet_utils import *
from scipy.stats import randint
from scipy.stats import uniform

project_dir = f"{os.getcwd()}.
↳split('SIDS_revelation_project')[0]}SIDS_revelation_project/"
image_dataset_path = f"{project_dir}datasets/onback_onstomach_v3"
model_path_fd = f"{project_dir}/models/4.fd_weights/best.pt"
model_path_pe = f"{project_dir}/models/2.pe_weights/best.pt"

emb_builder = EmbeddingBuilder(model_path_fd, image_dataset_path,
↳"load", weights_path_pe=model_path_pe)
embeddings = emb_builder.create_embedding(flags=True, positions=True,
↳positions_normalized=True, geometric_info=True, k_positions_normalized=True,
↳k_geometric_info=True)
```

The autoreload extension is already loaded. To reload it, use:

```
%reload_ext autoreload
```

Extracting dataset info from .coco.json

```

file:-----
Dataset contains 4158 valid samples, and labels are {'baby_on_back': 1,
'baby_on_stomach': 2}
-----

Loading features from
.csv-----
Features loaded succesfully, in particular there are 4158 files in the dataset
-----

Embedding builder initialized
successfully-----
Face detection model: 4 (YOLOv8)
Dataset: /Users/lorenzodimaio/Download/SIDS_revelation_project/datasets/onback_o
nstomach_v3
Dataset dimension: 4158
Dataset labels: {'baby_safe': 0, 'baby_unsafe': 1}
-----

Embedding
creation-----
Features: ['flag_eye1', 'flag_eye2', 'flag_nose', 'flag_mouth', 'x_eye1',
'y_eye1', 'x_eye2', 'y_eye2', 'x_nose', 'y_nose', 'x_mouth', 'y_mouth',
'x_eye1_norm', 'y_eye1_norm', 'x_eye2_norm', 'y_eye2_norm', 'x_nose_norm',
'y_nose_norm', 'x_mouth_norm', 'y_mouth_norm', 'eye_distance',
'eye_distance_norm', 'face_vertical_length', 'face_vertical_length_norm',
'face_angle_vertical', 'face_angle_horizontal', 'symmetry_diff', 'head_ration',
'x_nose_k', 'y_nose_k', 'x_left_eye_k', 'y_left_eye_k', 'x_right_eye_k',
'y_right_eye_k', 'x_left_ear', 'y_left_ear', 'x_right_ear', 'y_right_ear',
'x_left_shoulder', 'y_left_shoulder', 'x_right_shoulder', 'y_right_shoulder',
'x_left_elbow', 'y_left_elbow', 'x_right_elbow', 'y_right_elbow',
'x_left_wrist', 'y_left_wrist', 'x_right_wrist', 'y_right_wrist', 'x_left_hip',
'y_left_hip', 'x_right_hip', 'y_right_hip', 'x_left_knee', 'y_left_knee',
'x_right_knee', 'y_right_knee', 'x_left_ankle', 'y_left_ankle', 'x_right_ankle',
'y_right_ankle', 'shoulders_dist', 'shoulder_hip_right_dist',
'shoulder_hip_left_dist', 'nose_shoulder_right', 'nose_shoulder_left',
'shoulder_left_knee_right', 'shoulder_right_knee_left', 'knee_ankle_right',
'knee_ankle_left', 'nose_hip_right', 'nose_hip_left',
'elbow_shoulder_hip_right', 'elbow_shoulder_hip_left',
'shoulder_elbow_wrist_right', 'shoulder_elbow_wrist_left',
'shoulder_hip_knee_right', 'shoulder_hip_knee_left', 'hip_knee_ankle_right',
'hip_knee_ankle_left', 'shoulders_line_inclination', 'hips_line_inclination',
'torsion']
FINISHED: 4158 embedding created
-----

```

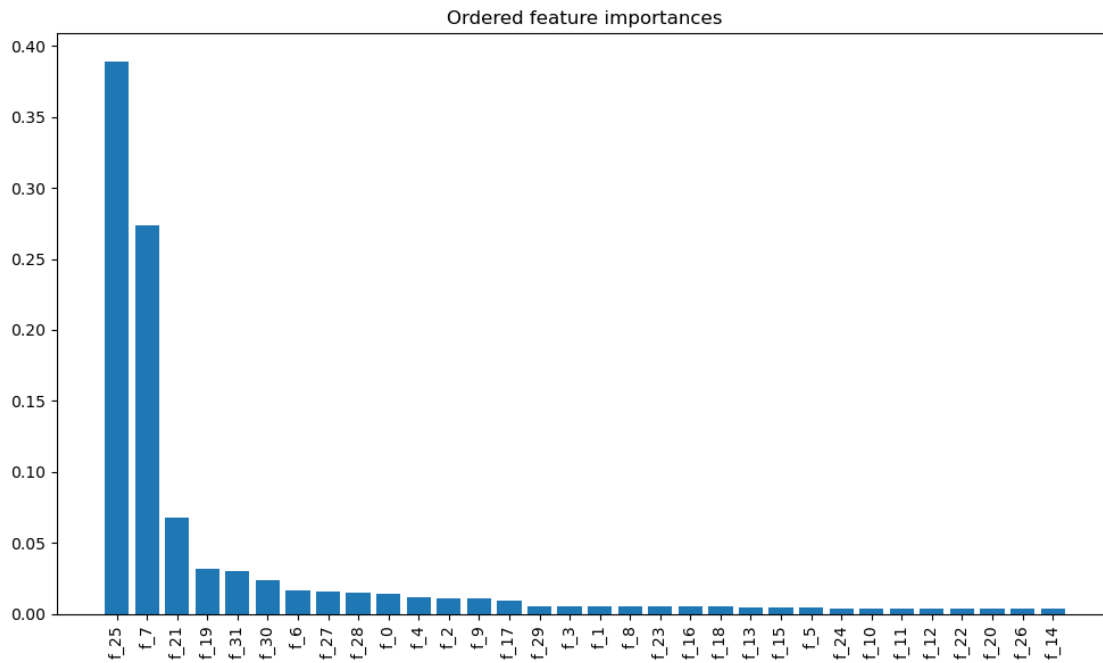
```


[19]: dataset = EmbeddingDataset(embeddings.to_numpy(), emb_builder.y, device=device)
      model = dataset.train_embeddings(embed_dim=32, epochs=50, batch_size=128,
      ↪lr=1e-3, verbose=False, weight_decay=1e-7, dropout_rate=0.05)

[20]: embeddings_new = dataset.extract_embeddings(model)
      embeddings_new= pd.DataFrame(embeddings_new.to_numpy(), columns=[f"f_{i}" for i
      ↪in range(embeddings_new.shape[1])])
      clf = Classifier(embeddings_new, dataset.y, emb_builder.classes_bs, figsize =
      ↪(10, 6))
      xgbc_1, results_xgbc_1 = clf.XGBC(verbose=True)

```

-----FIRST
ANALYSIS-----






Dataset labels:-----
 {'baby_safe': 0, 'baby_unsafe': 1}

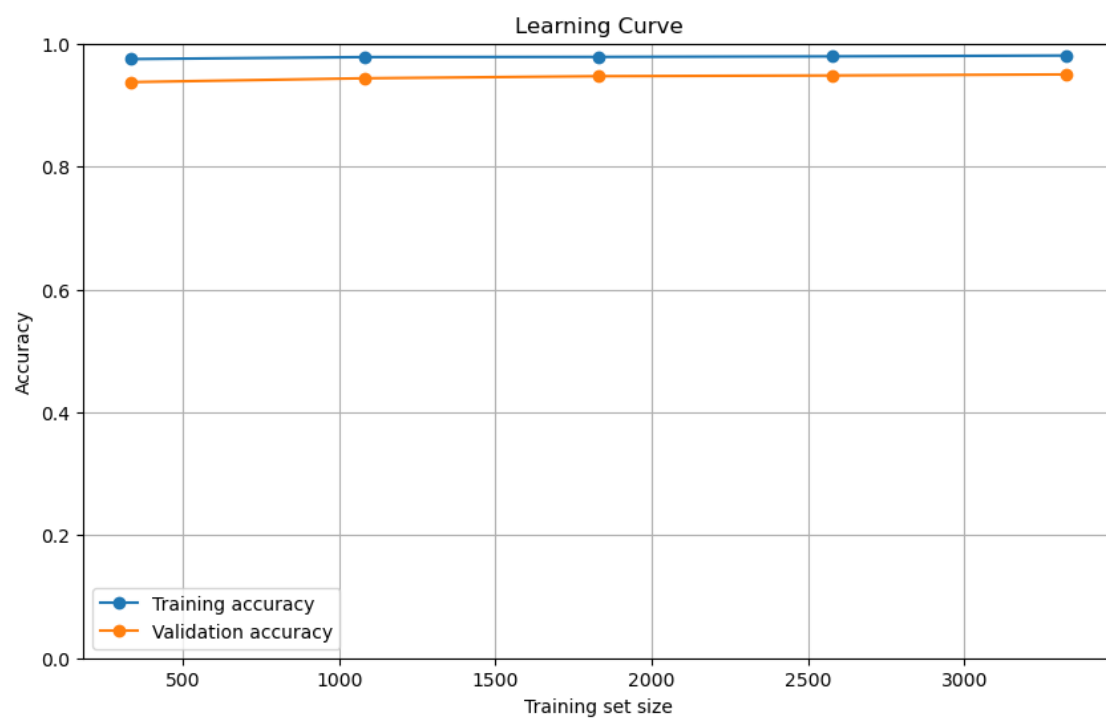
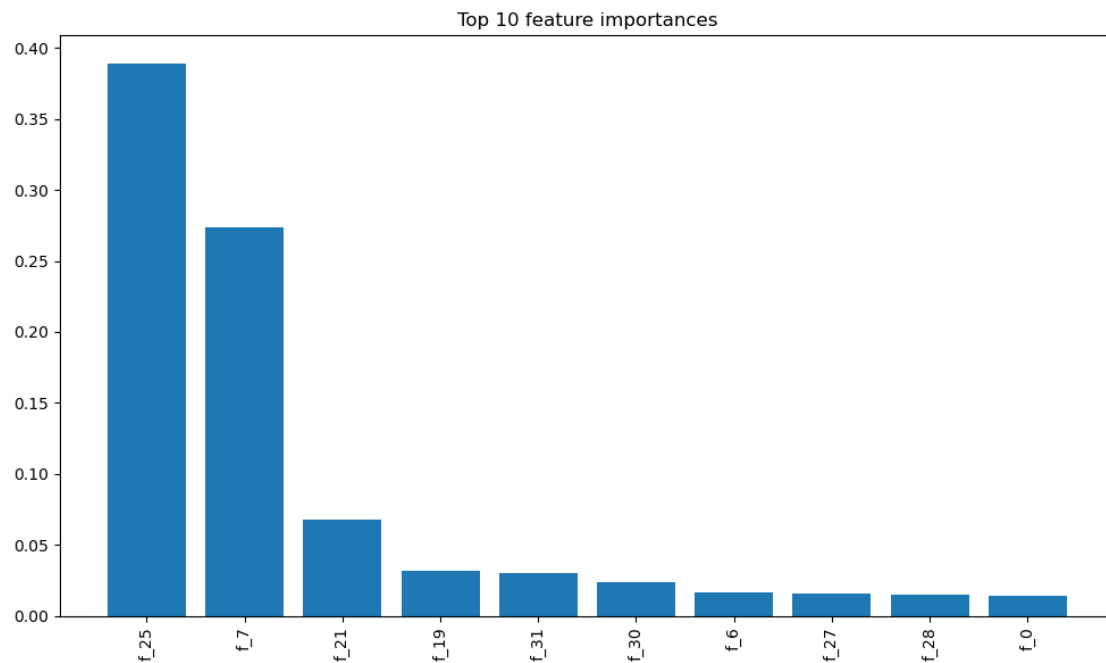
Report-----

	precision	recall	f1-score	support
baby_safe	0.95	0.95	0.95	423
baby_unsafe	0.94	0.95	0.95	409
accuracy			0.95	832
macro avg	0.95	0.95	0.95	832
weighted avg	0.95	0.95	0.95	832

Confusion matrix-----



-----TOP 10 FEATURES
ANALYSIS-----




Dataset labels:-----
 {'baby_safe': 0, 'baby_unsafe': 1}

Report-----



	precision	recall	f1-score	support
baby_safe	0.95	0.94	0.95	423
baby_unsafe	0.94	0.95	0.95	409
accuracy			0.95	832
macro avg	0.95	0.95	0.95	832
weighted avg	0.95	0.95	0.95	832

Confusion matrix-----



-----TOP 25 FEATURES

ANALYSIS-----




Dataset labels:-----
 {'baby_safe': 0, 'baby_unsafe': 1}

Report-----

	precision	recall	f1-score	support
baby_safe	0.96	0.94	0.95	423
baby_unsafe	0.94	0.96	0.95	409
accuracy			0.95	832
macro avg	0.95	0.95	0.95	832
weighted avg	0.95	0.95	0.95	832

Confusion matrix-----



```
[21]: clf.evaluation_pipeline_with_cv(xgbc_1,verbose=True)
```


-----K-FOLD CROSS-VALIDATION
ANALYSIS-----

Accuracy for each fold: [0.9519230769230769, 0.9555288461538461,
0.9567307692307693, 0.9446450060168472, 0.9530685920577617]
Average cross-validation accuracy: 0.9524

```
[21]: np.float64(0.9523792580764603)
```

1 1) Ottimizziamo i parametri della rete

```
[22]: import itertools

param_grid = {
    "embed_dim": [32, 64],
    "batch_size": [64, 128],
    "lr": [1e-3, 1e-4],
    "epochs": [30, 50, 100],
    "weight_decay": [1e-5, 1e-6, 1e-7],
    "dropout_rate": [0.0, 0.1, 0.2, 0.3]
}

best_acc = 0
best_params = None

for embed_dim, batch_size, lr, epochs, weight_decay, dropout_rate in itertools.
    product(
        param_grid["embed_dim"],
        param_grid["batch_size"],
        param_grid["lr"],
        param_grid["epochs"],
        param_grid["weight_decay"],
        param_grid["dropout_rate"]):

    # La tua istanziamento del modello e il training
    model = EmbeddingNet(input_dim=dataset.X.shape[1], embed_dim=embed_dim,
        dropout_rate=dropout_rate).to(device)

    trained_model, val_acc = dataset.train_and_evaluate(
        model=model,
        batch_size=batch_size,
        lr=lr,
        epochs=epochs,
        weight_decay=weight_decay,
        verbose=False
    )

    if val_acc > best_acc:
        best_acc = val_acc
        # Correzione: salva tutti gli iperparametri
```

```

best_params = {
    "embed_dim": embed_dim,
    "batch_size": batch_size,
    "lr": lr,
    "epochs": epochs,
    "weight_decay": weight_decay,
    "dropout_rate": dropout_rate
}

print("Best validation accuracy:", best_acc)
print("Best hyperparameters:", best_params)

```

Best validation accuracy: 0.9001203369434416

Best hyperparameters: {'embed_dim': 64, 'batch_size': 64, 'lr': 0.001, 'epochs': 100, 'weight_decay': 1e-07, 'dropout_rate': 0.3}

```

[23]: model = dataset.train_embeddings(**best_params, verbose=False)
      embeddings_new = dataset.extract_embeddings(model)
      embeddings_new = pd.DataFrame(embeddings_new.to_numpy(), columns=[f"f_{i}" for i
      ↪ in range(embeddings_new.shape[1])])


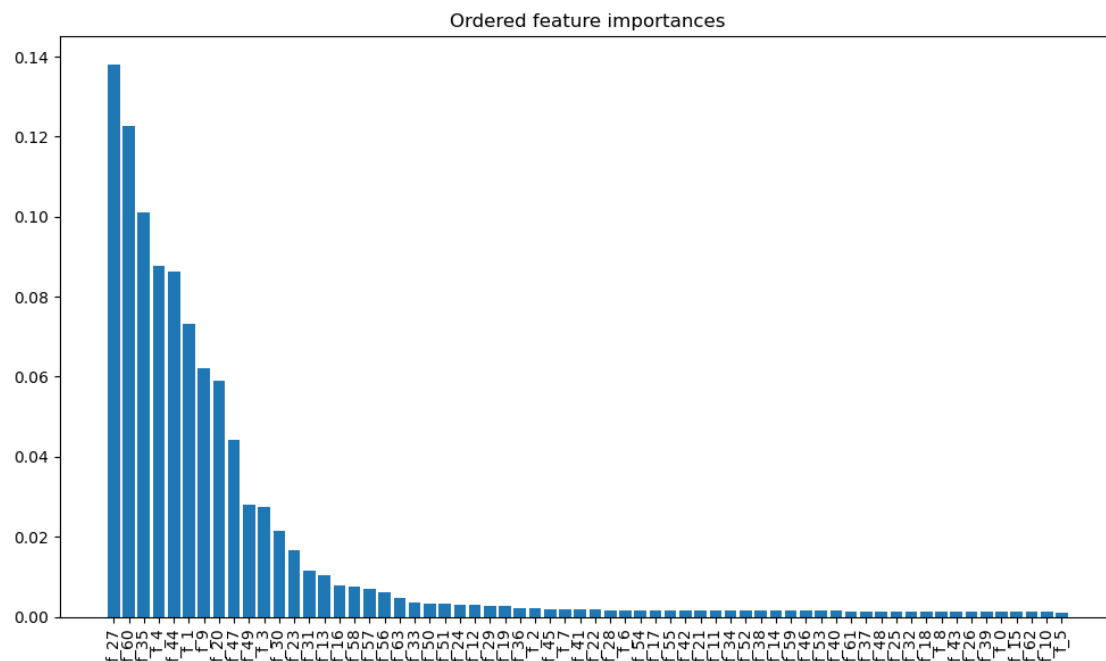
      clf = Classifier(embeddings_new, dataset.y, emb_builder.classes_bs, figsize =
      ↪ (10, 6))
      xgbc_1, results_xgbc_1 = clf.XGBC(verbose=True)

```

```

-----
-----
-----FIRST
ANALYSIS-----
-----
-----

```




Dataset labels:-----
 {'baby_safe': 0, 'baby_unsafe': 1}

Report-----


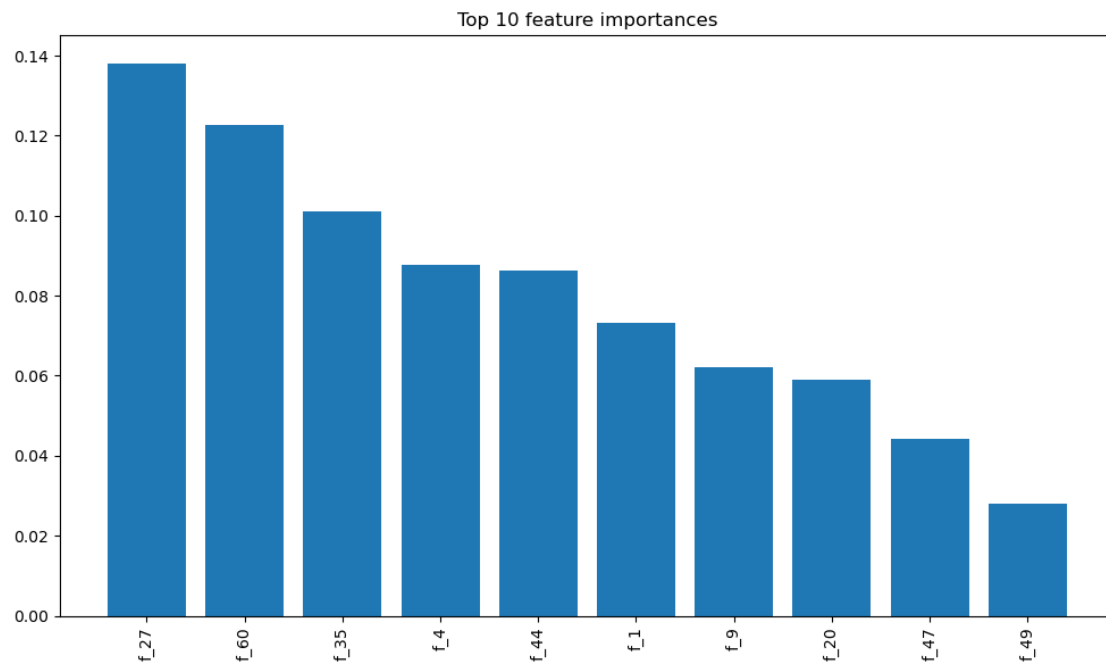
	precision	recall	f1-score	support
baby_safe	0.95	0.94	0.95	423
baby_unsafe	0.94	0.95	0.95	409
accuracy			0.95	832
macro avg	0.95	0.95	0.95	832
weighted avg	0.95	0.95	0.95	832

Confusion matrix-----



-----TOP 10 FEATURES

ANALYSIS-----




Dataset labels:-----
{'baby_safe': 0, 'baby_unsafe': 1}



Report-----

	precision	recall	f1-score	support
baby_safe	0.96	0.93	0.94	423
baby_unsafe	0.93	0.96	0.94	409
accuracy			0.94	832
macro avg	0.94	0.94	0.94	832
weighted avg	0.94	0.94	0.94	832

Confusion matrix-----



-----TOP 25 FEATURES
ANALYSIS-----




Dataset labels:-----
 {'baby_safe': 0, 'baby_unsafe': 1}

Report-----

	precision	recall	f1-score	support
baby_safe	0.96	0.94	0.95	423
baby_unsafe	0.94	0.96	0.95	409
accuracy			0.95	832
macro avg	0.95	0.95	0.95	832
weighted avg	0.95	0.95	0.95	832

Confusion matrix-----



```
[24]: clf.evaluation_pipeline_with_cv(xgbc_1,verbose=True)
```


-----K-FOLD CROSS-VALIDATION
ANALYSIS-----

Accuracy for each fold: [0.953125, 0.9567307692307693, 0.9543269230769231,
0.9434416365824309, 0.9446450060168472]
Average cross-validation accuracy: 0.9505

```
[24]: np.float64(0.9504538669813941)
```

2 2) Ottimizziamo i parametri del modello

```
[35]: from scipy.stats import randint
      from scipy.stats import uniform

      param_dist = {
          'n_estimators': randint(100, 500), # Numero di alberi da costruire
          'learning_rate': uniform(0.01, 0.3), # Passo di apprendimento
          'max_depth': randint(3, 10), # Profondità massima di un albero
          'subsample': uniform(0.6, 0.4), # Frazione di campioni da usare per
          ↳ l'addestramento
          'colsample_bytree': uniform(0.6, 0.4), # Frazione di feature da usare per
          ↳ la costruzione di ogni albero
          'gamma': uniform(0, 0.5) # Riduzione minima della perdita richiesta per
          ↳ fare un'ulteriore partizione
      }
      best_params_xgbc = clf.optimize_model(xgbc_1, param_dist, verbose=True)
      xgbc_2 = XGBClassifier(**best_params_xgbc)
      results_xgbc_2 = clf.evaluation_pipeline(xgbc_2, verbose=True, optimized=True)
```


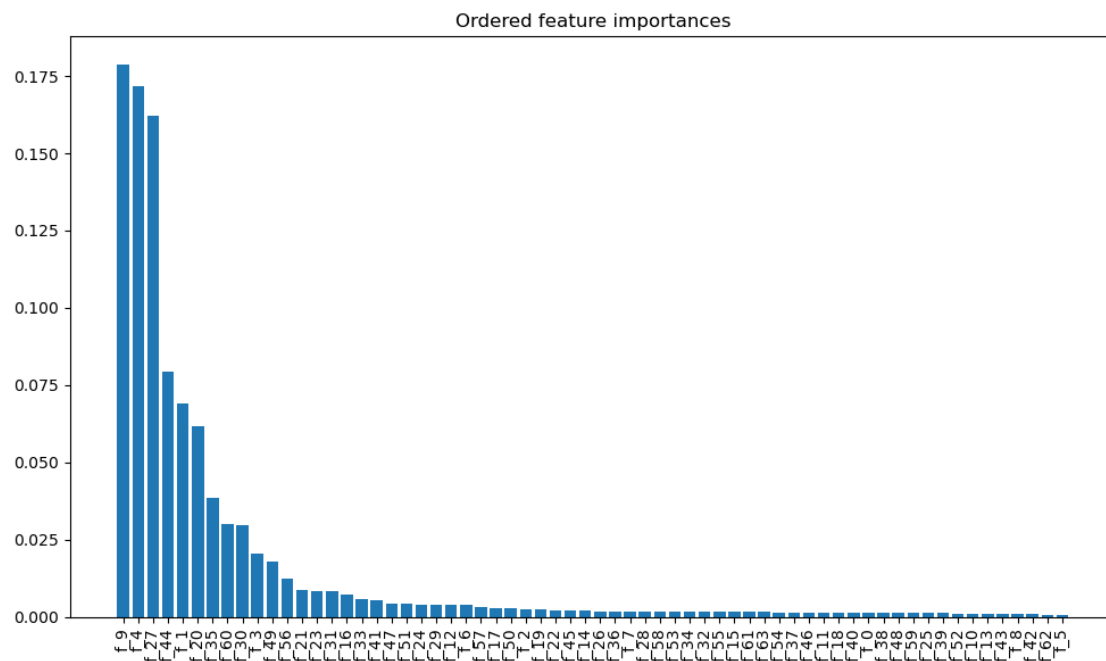
Start random search...

Random Search Results:

Best parameters found: {'colsample_bytree': np.float64(0.7970472375971478),
'gamma': np.float64(0.1643758051437541), 'learning_rate':
np.float64(0.20002025629501774), 'max_depth': 5, 'n_estimators': 470,
'subsample': np.float64(0.8149053666856281)}
Best mean cross-validation accuracy: 0.9506929485876855

Evaluation of the best model on the test set:

```
-----
-----
-----FIRST
ANALYSIS-----
-----
-----
```




Dataset labels:-----
{'baby_safe': 0, 'baby_unsafe': 1}

Report-----

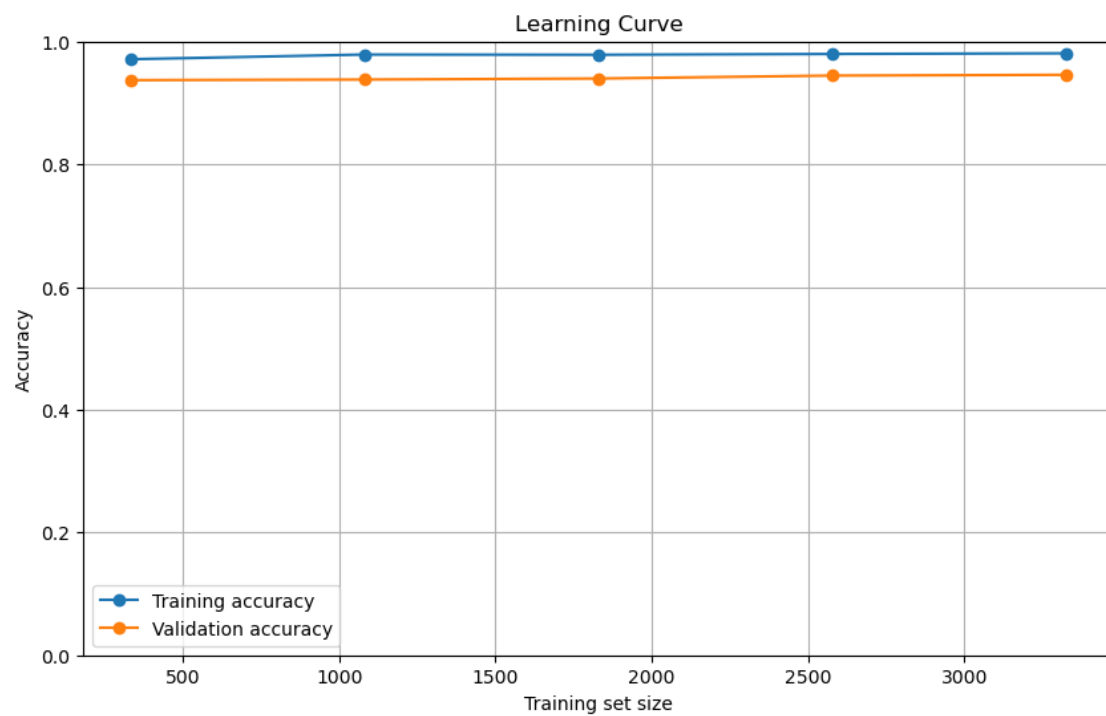
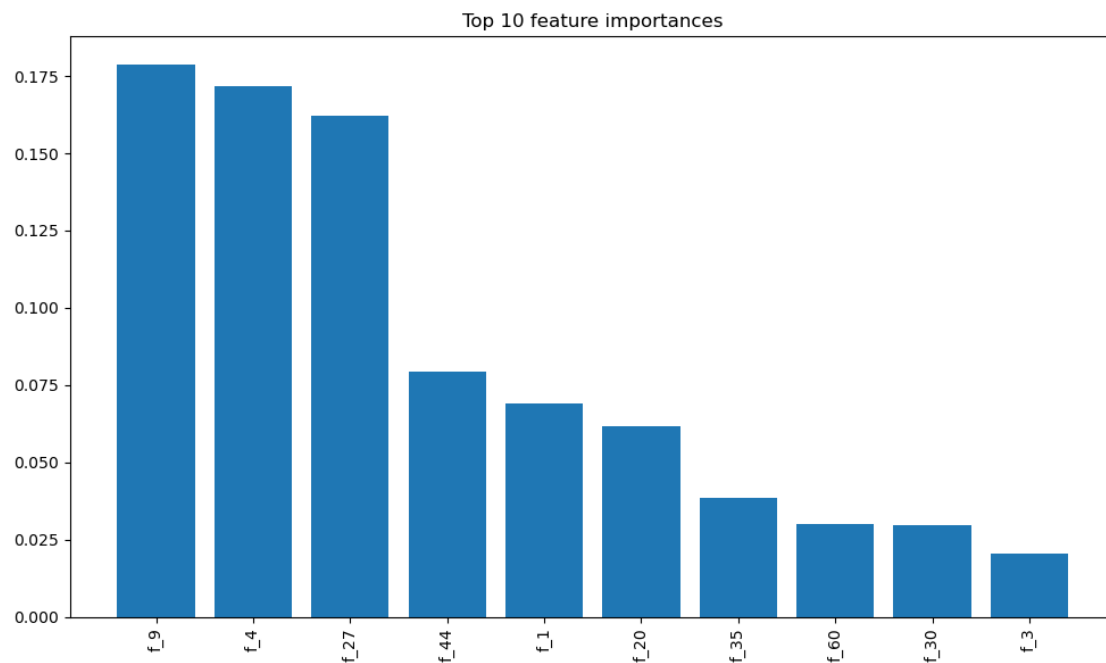
	precision	recall	f1-score	support
baby_safe	0.95	0.94	0.94	423
baby_unsafe	0.94	0.95	0.94	409
accuracy			0.94	832
macro avg	0.94	0.94	0.94	832
weighted avg	0.94	0.94	0.94	832

Confusion matrix-----



-----TOP 10 FEATURES

ANALYSIS-----




Dataset labels:-----
 {'baby_safe': 0, 'baby_unsafe': 1}

Report-----


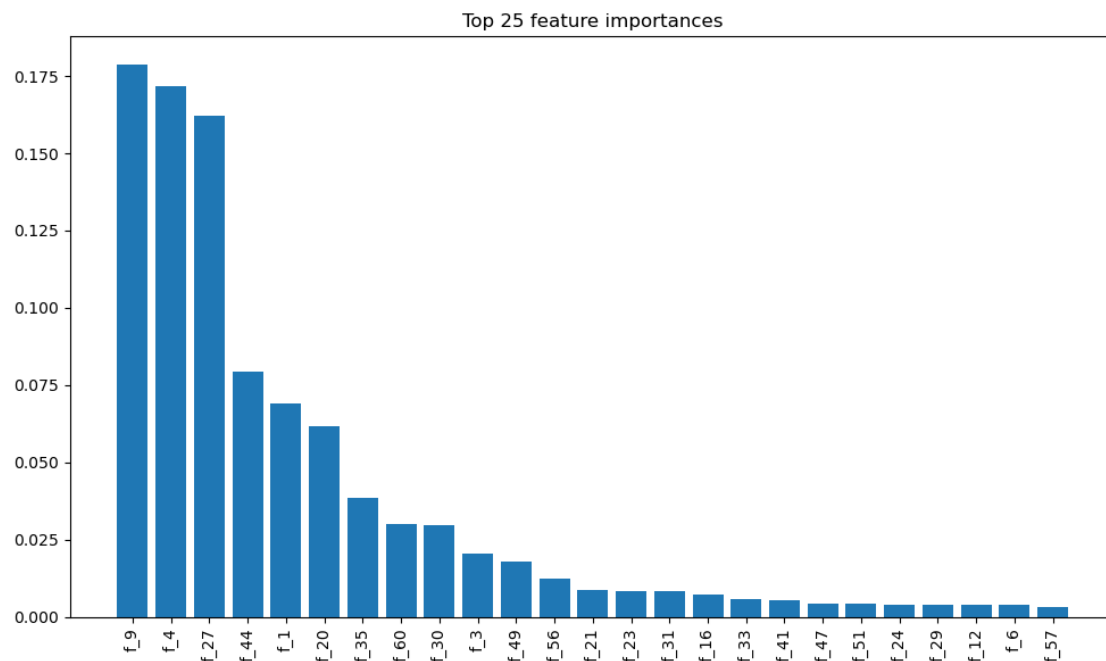
	precision	recall	f1-score	support
baby_safe	0.96	0.93	0.94	423
baby_unsafe	0.93	0.96	0.94	409
accuracy			0.94	832
macro avg	0.94	0.94	0.94	832
weighted avg	0.94	0.94	0.94	832

Confusion matrix-----



-----TOP 25 FEATURES

ANALYSIS-----




Dataset labels:-----
 {'baby_safe': 0, 'baby_unsafe': 1}

Report-----

	precision	recall	f1-score	support
baby_safe	0.95	0.94	0.95	423
baby_unsafe	0.94	0.95	0.94	409
accuracy			0.94	832
macro avg	0.94	0.94	0.94	832
weighted avg	0.94	0.94	0.94	832

Confusion matrix-----



```
[37]: dataset.save_embeddings(embeddings_new.iloc[:, :10])
```

```
[32]: xgbc_3 = XGBClassifier(**best_params_xgbc)
      clf.evaluation_pipeline_with_cv(xgbc_3, verbose=True)
```


-----K-FOLD CROSS-VALIDATION

ANALYSIS-----

Accuracy for each fold: [0.953125, 0.9555288461538461, 0.9507211538461539,
0.941034897713598, 0.9446450060168472]
Average cross-validation accuracy: 0.9490

[32]: np.float64(0.9490109807460889)

3 3) Eseguiamo Ablazione

```
[38]: emb_reduced, history = clf.iterative_ablation(  
    model=xgbc_2,  
    embeddings=embeddings_new,  
    y=dataset.y,  
    classes_bs=emb_builder.classes_bs,  
    max_cycles=20,  
    n_very_big_groups=10,  
    very_big_group_size=10,  
    n_big_groups=10,  
    big_group_size=6,  
    n_medium_groups=3,  
    medium_group_size=10,  
    n_small_groups=10,  
    small_group_size=2,  
    verbose=False)
```



```
[39]: diff1 = set(embeddings_new.columns) - set(emb_reduced.columns) # colonne in  
    ↪ df1 ma non in df2  
  
print("Colonne presenti in df1 ma non in df2:", diff1)
```

Colonne presenti in df1 ma non in df2: {'f_8'}

```
[40]: clf = Classifier(emb_reduced, dataset.y, emb_builder.classes_bs)  
xgbc_4 = XGBClassifier(**best_params_xgbc)  
clf.evaluation_pipeline(xgbc_4, verbose=True, optimized=True)
```

-----FIRST

ANALYSIS-----




Dataset labels:-----
{'baby_safe': 0, 'baby_unsafe': 1}


Report-----


	precision	recall	f1-score	support
baby_safe	0.96	0.95	0.95	423
baby_unsafe	0.94	0.96	0.95	409
accuracy			0.95	832
macro avg	0.95	0.95	0.95	832
weighted avg	0.95	0.95	0.95	832

Confusion matrix-----



-----TOP 10 FEATURES
ANALYSIS-----






Dataset labels:-----
{'baby_safe': 0, 'baby_unsafe': 1}



Report-----

	precision	recall	f1-score	support
baby_safe	0.96	0.93	0.95	423
baby_unsafe	0.93	0.96	0.95	409
accuracy			0.95	832
macro avg	0.95	0.95	0.95	832
weighted avg	0.95	0.95	0.95	832

Confusion matrix-----



-----TOP 25 FEATURES
ANALYSIS-----




Dataset labels:-----
{'baby_safe': 0, 'baby_unsafe': 1}

Report-----

	precision	recall	f1-score	support
baby_safe	0.95	0.94	0.95	423
baby_unsafe	0.94	0.95	0.94	409
accuracy			0.94	832
macro avg	0.94	0.94	0.94	832
weighted avg	0.94	0.94	0.94	832

Confusion matrix-----




```

[40]: {'all_features': {'model': XGBClassifier(base_score=None, booster=None,
callbacks=None,
        colsample_bylevel=None, colsample_bynode=None,
        colsample_bytree=np.float64(0.7970472375971478), device=None,
        early_stopping_rounds=None, enable_categorical=False,
        eval_metric=None, feature_types=None, feature_weights=None,
        gamma=np.float64(0.1643758051437541), grow_policy=None,
        importance_type=None, interaction_constraints=None,
        learning_rate=np.float64(0.20002025629501774), max_bin=None,
        max_cat_threshold=None, max_cat_to_onehot=None,
        max_delta_step=None, max_depth=5, max_leaves=None,
        min_child_weight=None, missing=nan, monotone_constraints=None,
        multi_strategy=None, n_estimators=470, n_jobs=None,
        num_parallel_tree=None, ...),
        'X': array([[ -0.1247,   0.012603,   0.14913, ...,  -0.22047,
0.082049,   0.15007],
        [ -0.081928,  -0.14025,   0.15321, ...,  -0.21451,   0.087326,
-0.019009],
        [ -0.14457,  -0.029547,   0.15554, ...,  -0.23851,   0.082482,
0.11234],
        ...,
        [ -0.14506,  -0.0065284,   0.15285, ...,  -0.23238,   0.087467,
0.13656],
        [ -0.13297,   0.0034594,   0.13928, ...,  -0.20727,   0.081965,
0.16009],
        [ -0.11506,  -0.031399,   0.16687, ...,  -0.22957,   0.098249,
0.11729]], shape=(4158, 63), dtype=float32),
        'y': array([0, 1, 0, ..., 0, 0, 0], shape=(4158,))},
        'top_10_features': {'model': XGBClassifier(base_score=None, booster=None,
callbacks=None,
        colsample_bylevel=None, colsample_bynode=None,
        colsample_bytree=np.float64(0.7970472375971478), device=None,
        early_stopping_rounds=None, enable_categorical=False,
        eval_metric=None, feature_types=None, feature_weights=None,
        gamma=np.float64(0.1643758051437541), grow_policy=None,
        importance_type=None, interaction_constraints=None,
        learning_rate=np.float64(0.20002025629501774), max_bin=None,
        max_cat_threshold=None, max_cat_to_onehot=None,
        max_delta_step=None, max_depth=5, max_leaves=None,
        min_child_weight=None, missing=nan, monotone_constraints=None,
        multi_strategy=None, n_estimators=470, n_jobs=None,
        num_parallel_tree=None, ...),
        'X': array([[ -0.11771,  -0.068812,   0.030948, ...,   0.088932,
0.13121,  -0.12217],
        [ 0.0016136,   0.070188,   0.13179, ...,  -0.10022,   0.03058,
0.076435],
        [ -0.10086,  -0.037679,   0.061949, ...,   0.055472,   0.11399,

```

```

-0.09763],
    ...,
    [ -0.12222, -0.045986, 0.043369, ..., 0.074299, 0.11156,
-0.12396],
    [ -0.13586, -0.085682, 0.012979, ..., 0.10869, 0.12254,
-0.13611],
    [ -0.1145, -0.028314, 0.058142, ..., 0.056774, 0.12513,
-0.077972]], shape=(4158, 10), dtype=float32),
    'y': array([0, 1, 0, ..., 0, 0, 0], shape=(4158,)),
    'top_features_idx': array([26, 34, 8, 43, 40, 19, 59, 48, 3, 22])},
    'top_25_features': {'model': XGBClassifier(base_score=None, booster=None,
callbacks=None,
                    colsample_bylevel=None, colsample_bynode=None,
                    colsample_bytree=np.float64(0.7970472375971478), device=None,
                    early_stopping_rounds=None, enable_categorical=False,
                    eval_metric=None, feature_types=None, feature_weights=None,
                    gamma=np.float64(0.1643758051437541), grow_policy=None,
                    importance_type=None, interaction_constraints=None,
                    learning_rate=np.float64(0.20002025629501774), max_bin=None,
                    max_cat_threshold=None, max_cat_to_onehot=None,
                    max_delta_step=None, max_depth=5, max_leaves=None,
                    min_child_weight=None, missing=nan, monotone_constraints=None,
                    multi_strategy=None, n_estimators=470, n_jobs=None,
                    num_parallel_tree=None, ...),
    'X': array([[ -0.11771, -0.068812, 0.030948, ..., -0.066666,
-0.20229, 0.14913],
    [ 0.0016136, 0.070188, 0.13179, ..., -0.12813, -0.15578,
0.15321],
    [ -0.10086, -0.037679, 0.061949, ..., -0.095126, -0.18774,
0.15554],
    ...,
    [ -0.12222, -0.045986, 0.043369, ..., -0.083096, -0.19705,
0.15285],
    [ -0.13586, -0.085682, 0.012979, ..., -0.068441, -0.20908,
0.13928],
    [ -0.1145, -0.028314, 0.058142, ..., -0.084555, -0.2168,
0.16687]], shape=(4158, 25), dtype=float32),
    'y': array([0, 1, 0, ..., 0, 0, 0], shape=(4158,)),
    'top_features_idx': array([26, 34, 8, 43, 40, 19, 59, 48, 3, 22, 4, 15, 30,
46, 57, 12, 6, 1, 50, 11, 23, 32, 18, 36, 2])}}

```

```
[41]: clf.evaluation_pipeline_with_cv(xgbc_4,verbose=True)
```

```

-----
-----
-----K-FOLD CROSS-VALIDATION
ANALYSIS-----
-----

```

```
-----  
Accuracy for each fold: [0.953125, 0.9579326923076923, 0.9495192307692307,  
0.9398315282791817, 0.941034897713598]  
Average cross-validation accuracy: 0.9483
```

```
[41]: np.float64(0.9482886698139404)
```

```
[112]: os.environ["PATH"] = "/Library/TeX/texbin:" + os.environ["PATH"]  
save_as_pdf(ipynbname.path())
```

```
[NbConvertApp] Converting notebook /Users/lorenzodimaio/Download/SIDS_revelation_project/full_pipeline/reports/MLP_classifier_parameter.ipynb to pdf  
[NbConvertApp] Support files will be in /Users/lorenzodimaio/Download/SIDS_revelation_project/full_pipeline/reports/reports/MLP_classifier_parameter(2025-09-09)_files/  
[NbConvertApp] Writing 70178 bytes to notebook.tex  
[NbConvertApp] Building PDF  
[NbConvertApp] Running xelatex 3 times: ['xelatex', 'notebook.tex', '-quiet']  
[NbConvertApp] Running bibtex 1 time: ['bibtex', 'notebook']  
[NbConvertApp] WARNING | bibtex had problems, most likely because there were no citations  
[NbConvertApp] PDF successfully created  
[NbConvertApp] Writing 831636 bytes to /Users/lorenzodimaio/Download/SIDS_revelation_project/full_pipeline/reports/reports/MLP_classifier_parameter(2025-09-09).pdf
```